

EXTENDS *Apache*, *U2PC*, *TLC*

@type:  $(a, b) \Rightarrow \langle a, b \rangle$ ;  
 $Pair(A, B) \triangleq \langle A, B \rangle$

1 shard, 1-2 transactions

Checking simple commit, and conflict behaviours

$T1 \triangleq SetAsFun(\{Pair("T1", \{ "X" \})\})$   
 $T1\_2 \triangleq SetAsFun(\{Pair("T1", \{ "X" \}), Pair("T2", \{ "X" \})\})$   
 $S1 \triangleq SetAsFun(\{Pair("X", \{ "X1", "X2" \})\})$

3 shards, 3 transactions

Checking indirect dependency loops

$T3 \triangleq SetAsFun(\{$   
 $Pair("T1", \{ "X", "Y" \}),$   
 $Pair("T2", \{ "Y", "Z" \}),$   
 $Pair("T3", \{ "Z", "X" \})\})$   
 $S3 \triangleq SetAsFun(\{$   
 $Pair("X", \{ "X1", "X2" \}),$   
 $Pair("Y", \{ "Y1", "Y2" \}),$   
 $Pair("Z", \{ "Z1", "Z2" \})\})$

Initial state for *Apache* testing

$CInit \triangleq$   
 $\wedge Txns := T3$   
 $\wedge Shards := S3$

Credit to <https://github.com/tlaplus/examples>

$TransitiveClosure(R) \triangleq$   
 $LET\ S \triangleq \{r[1] : r \in R\} \cup \{r[2] : r \in R\}$   
 $RECURSIVE\ TCR(-)$   
 $TCR(T) \triangleq IF\ T = \{\}$   
 $THEN\ R$   
 $ELSE\ LET\ r \triangleq CHOOSE\ s \in T : TRUE$   
 $RR \triangleq TCR(T \setminus \{r\})$   
 $IN\ RR \cup \{\langle s, t \rangle \in S \times S :$   
 $\langle s, r \rangle \in RR \wedge \langle r, t \rangle \in RR\}$   
 $IN\ TCR(S)$

$TransactionOrdering \triangleq LET$   
 $F(acc, tid) \triangleq acc \cup (Range(Coordinator\_txn\_state[tid]) \times \{tid\})$   
 $Base \triangleq ApaFoldSet(F, \{\}, TIDs)$   
 $IN\ TransitiveClosure(Base)$

$$\begin{aligned}
\text{RecoveryCommitted}(S) &\triangleq \\
&\{t \in \text{TIDs} : \\
&\quad \forall r \in S : \\
&\quad \quad \text{KeyLookup}[r] \in \text{Trns}[t] \\
&\quad \quad \Rightarrow \vee \text{Replicas}[r].\text{locked} \wedge \text{Replicas}[r].\text{logged} = t \\
&\quad \quad \vee \text{Replicas}[r].\text{version} = t \\
&\quad \quad \vee \langle t, \text{Replicas}[r].\text{version} \rangle \in \text{TransactionOrdering} \\
&\}
\end{aligned}$$

Every transaction committed during recovery preserves linearisability

$$\begin{aligned}
\text{Safety\_recovery} &\triangleq \\
&\forall S \in \text{SUBSET RIDs} : \\
&\quad \text{Valid recovery} \\
&\quad (\forall k \in \text{DOMAIN Shards} : \exists r \in S : r \in \text{Shards}[k]) \\
&\quad \Rightarrow \text{Linearisability}(\text{CommittedTIDs} \cup \text{RecoveryCommitted}(S))
\end{aligned}$$

$$\begin{aligned}
\text{RecoveryAborted}(S) &\triangleq \\
&\{t \in \text{TIDs} : \\
&\quad \exists r \in S : \\
&\quad \quad \wedge \text{KeyLookup}[r] \in \text{Trns}[t] \\
&\quad \quad \wedge \vee \neg \text{Replicas}[r].\text{locked} \\
&\quad \quad \vee \text{Replicas}[r].\text{locked} \wedge \text{Replicas}[r].\text{logged} \neq t\}
\end{aligned}$$

Every committed or aborted transaction results in the same recovery decision

$$\begin{aligned}
\text{Durability} &\triangleq \\
&\forall S \in \text{SUBSET RIDs} : \\
&\quad \text{Valid recovery} \\
&\quad (\forall k \in \text{DOMAIN Shards} : \exists r \in S : r \in \text{Shards}[k]) \\
&\quad \Rightarrow \\
&\quad \forall t \in \text{TIDs} : \\
&\quad \quad \wedge t \in \text{CommittedTIDs} \Rightarrow t \in \text{RecoveryCommitted}(S) \\
&\quad \quad \wedge t \in \text{AbortedTIDs} \Rightarrow t \in \text{RecoveryAborted}(S)
\end{aligned}$$

Since recovery stops every replica it uses, an explicit recovery check is unnecessary since that is equivalent to just checking that every possible recovery using the current state preserves the invariants.

$$\begin{aligned}
\text{Invs} &\triangleq \\
&\quad \wedge \text{Safety\_recovery} \\
&\quad \wedge \text{Durability}
\end{aligned}$$